

What is Pandas?

- A Python library built on top of NumPy.
- Designed for working with **tabular data** (rows & columns).
- Widely used in **finance, data science, and statistics**.
- Two main objects:
 - **Series** → 1D labeled array (like one column in Excel).
 - **DataFrame** → 2D labeled data structure (like an Excel sheet).

NumPy vs Pandas

Feature	NumPy Array	Pandas DataFrame
Type of data	Homogeneous (all same type)	Heterogeneous (different types)
Labels	Indexed only by numbers	Columns and row labels
Structure	Mostly numeric	Tabular (rows + columns)
Use case	Mathematical operations	Data analysis & manipulation

Creating a Pandas Series

python

```
import pandas as pd

# From a Python List
s = pd.Series([100, 102, 98, 105, 110], name="Stock Prices")
print(s)
```

Output:

yaml

```
0    100
1    102
2     98
3    105
4    110
Name: Stock Prices, dtype: int64
```

👉 Explanation:

- The left column is the **index** (default: 0,1,2...).
- The right column is the **values**.
- The `name` attribute labels the Series.

Creating a Pandas DataFrame

python

```
data = {  
    "Stock": ["AAPL", "GOOG", "AMZN", "MSFT", "TSLA"],  
    "Price": [150, 2800, 3400, 299, 720],  
    "Volume": [100000, 150000, 120000, 90000, 110000]  
}  
  
df = pd.DataFrame(data)  
print(df)
```

Output:

yaml

	Stock	Price	Volume
0	AAPL	150	100000
1	GOOG	2800	150000
2	AMZN	3400	120000
3	MSFT	299	90000
4	TSLA	720	110000

👉 Explanation:

- Each **column has a name** ("Stock", "Price", "Volume").
- Each **row has an index** (0 → 4).
- Looks similar to an Excel sheet.



Exercise – Pandas Basics with Series & DataFrame

1. Create a **Pandas Series** of 5 daily stock returns: `[0.01, -0.005, 0.02, -0.01, 0.015]`.
 - Each value corresponds to the **return of the stock on that day**.
 - Example: Day 1 → +1%, Day 2 → -0.5%, etc.
2. Create a **DataFrame** with two columns:
 - `"Day"` → numbers from 1 to 5
 - `"Return"` → the Series above
3. Print both the **Series** and the **DataFrame**.
4. Add a new column `"Cumulative Return"` to the DataFrame that shows the **cumulative product of returns** over time.
 - This shows the total growth of the investment if you reinvest profits each day.

Hint:

- You can use the Pandas function `.cumprod()` to calculate the cumulative product.
- Since returns are percentages, remember to add 1 before applying `.cumprod()`.
Example: `df["Cumulative Return"] = (1 + df["Return"]).cumprod()`

✓ Correction – Pandas Basics with Cumulative Return

```
pythonCopy  
  
import pandas as pd  
  
# 1. Create a Series of stock returns  
returns = pd.Series([0.01, -0.005, 0.02, -0.01, 0.015], name="Return")  
print("Series of Returns:")  
print(returns)  
  
# 2. Create a DataFrame with Day and Return  
data = {  
    "Day": [1, 2, 3, 4, 5],  
    "Return": returns  
}  
df = pd.DataFrame(data)  
  
print("\nDataFrame of Returns:")  
print(df)  
  
# 3. Add Cumulative Return column  
df["Cumulative Return"] = (1 + df["Return"]).cumprod()  
  
print("\nDataFrame with Cumulative Return:")  
print(df)
```

Series of Returns:

```
vbnet  
  
0      0.010  
1     -0.005  
2      0.020  
3     -0.010  
4      0.015  
Name: Return, dtype: float64
```

DataFrame with Cumulative Return:

```
sql  
  
      Day  Return  Cumulative Return  
0      1    0.010          1.010000  
1      2   -0.005          1.004950  
2      3    0.020          1.025049  
3      4   -0.010          1.014798  
4      5    0.015          1.030020
```

Topic 2: Indexing & Selection in Pandas

Why Indexing Matters in Pandas?

- The index allows you to access rows and columns efficiently.
- Think of it as the “address” of your data.
- Pandas offers two main methods:
 - `.loc[]` → label-based indexing (by names, not numbers).
 - `.iloc[]` → position-based indexing (by row/column number).

Example: Our DataFrame

python

```
import pandas as pd

data = {
    "Stock": ["AAPL", "GOOG", "AMZN", "MSFT", "TSLA"],
    "Price": [150, 2800, 3400, 299, 720],
    "Volume": [100000, 150000, 120000, 90000, 110000]
}
df = pd.DataFrame(data)
print(df)
```

Output:

yaml

	Stock	Price	Volume
0	AAPL	150	100000
1	GOOG	2800	150000
2	AMZN	3400	120000
3	MSFT	299	90000
4	TSLA	720	110000

Selecting Columns

python

```
print(df["Stock"])          # Single column  
print(df[["Stock", "Price"]]) # Multiple columns
```

Output:

yaml

```
0    AAPL  
1    GOOG  
2    AMZN  
3    MSFT  
4    TSLA  
Name: Stock, dtype: object
```

	Stock	Price
0	AAPL	150
1	GOOG	2800
2	AMZN	3400
3	MSFT	299
4	TSLA	720

Row Selection with loc (by labels)

python

```
print(df.loc[2])          # Row with index Label 2
print(df.loc[1:3])        # Rows from 1 to 3 (inclusive)
```

Output:

yaml

```
Stock      AMZN
Price      3400
Volume     120000
Name: 2, dtype: object
```

	Stock	Price	Volume
1	GOOG	2800	150000
2	AMZN	3400	120000
3	MSFT	299	90000

Row Selection with iloc (by position)

python

```
print(df.iloc[0])          # First row  
print(df.iloc[0:2])        # First two rows
```

Output:

yaml

```
Stock      AAPL  
Price      150  
Volume     100000  
Name: 0, dtype: object
```

```
   Stock  Price  Volume  
0  AAPL    150   100000  
1  GOOG    2800   150000
```

Boolean Indexing

python

```
print(df[df["Price"] > 1000])    # All stocks with Price > 1000
```

Output:

yaml

	Stock	Price	Volume
1	GOOG	2800	150000
2	AMZN	3400	120000

Mini-Exercise – Indexing with Returns DataFrame

Use the previous DataFrame of daily returns (with columns `Day`, `Return`, `Cumulative Return`) from the first exercise.

sql

 Copier le code

	Day	Return	Cumulative Return
0	1	0.010	1.010000
1	2	-0.005	1.004950
2	3	0.020	1.025049
3	4	-0.010	1.014798
4	5	0.015	1.030020

Tasks:

1. Select only the `"Return"` column.
2. Select the row corresponding to `Day = 3` using `.loc[]`.
3. Select the first 3 rows using `.iloc[]`.
4. Select all rows where `"Return"` is strictly positive (> 0).

1. Select only the "Return" column

```
returns_only = df["Return"]
```

```
print("Return column:")
```

```
print(returns_only)
```

2. Select the row corresponding to Day = 3 using loc

```
day3_row = df.loc[df["Day"] == 3]
```

```
print("\nRow where Day = 3:")
```

```
print(day3_row)
```

3. Select the first 3 rows using iloc

```
first_three_rows = df.iloc[0:3]
```

```
print("\nFirst 3 rows:")
```

```
print(first_three_rows)
```

4. Select all rows where Return > 0

```
positive_returns = df[df["Return"] > 0]
```

```
print("\nRows with positive Return:")
```

```
print(positive_returns)
```

Return column:

```
vbnet

0    0.010
1   -0.005
2    0.020
3   -0.010
4    0.015

Name: Return, dtype: float64
```

Row where Day = 3:

```
sql

    Day  Return  Cumulative Return
2     3   0.020             1.025049
```

First 3 rows:

```
sql

    Day  Return  Cumulative Return
0     1   0.010             1.010000
1     2  -0.005             1.004950
2     3   0.020             1.025049
```

Rows with positive Return:

```
sql

    Day  Return  Cumulative Return
0     1   0.010             1.010000
2     3   0.020             1.025049
4     5   0.015             1.030020
```

Topic 3: Operations on DataFrame

1. Adding and Modifying Columns

- You can create new columns based on existing data.
- Examples in finance: daily returns, cumulative returns, log returns, rolling volatility, etc.
- Syntax:

python

 Copier le code

```
df["NewColumn"] = ...
```

Example – Adding Rolling Volatility (3-day window)

python

 Copier le code

```
import pandas as pd
import numpy as np

# Assume df is the existing DataFrame with Day, Return, Cumulative Return
df["Rolling Volatility"] = df["Return"].rolling(window=3).std()
print(df)
```

Output:

pgsql

 Copier le code

	Day	Return	Cumulative Return	Rolling Volatility
0	1	0.010	1.010000	NaN
1	2	-0.005	1.004950	NaN
2	3	0.020	1.025049	0.012583
3	4	-0.010	1.014798	0.015275
4	5	0.015	1.030020	0.015275

💡 Note: First two rows are NaN because a 3-day window is required to calculate the standard deviation.

2. Using Built-in Aggregation Functions

- Pandas provides functions like `.mean()`, `.sum()`, `.min()`, `.max()`, `.std()` for quick statistics.

python

 Copy

```
mean_return = df["Return"].mean()
sum_return = df["Return"].sum()
min_return = df["Return"].min()
max_return = df["Return"].max()
std_return = df["Return"].std()

print(f"Mean Return: {mean_return}")
print(f"Sum of Returns: {sum_return}")
print(f"Min Return: {min_return}")
print(f"Max Return: {max_return}")
print(f"Std of Return: {std_return}")
```

Output:

mathematica

 Copy

```
Mean Return: 0.006
Sum of Returns: 0.03
Min Return: -0.01
Max Return: 0.02
Std of Return: 0.01224744871391589
```

3. Using `apply()` for Custom Operations

- Apply a function to a column (or row) using `.apply()`.
- Example: calculate squared returns:

python

```
df["Squared Return"] = df["Return"].apply(lambda x: x**2)
print(df)
```

Output:

pgsql

	Day	Return	Cumulative Return	Rolling Volatility	Squared Return
0	1	0.010	1.010000	NaN	0.000100
1	2	-0.005	1.004950	NaN	0.000025
2	3	0.020	1.025049	0.012583	0.000400
3	4	-0.010	1.014798	0.015275	0.000100
4	5	0.015	1.030020	0.015275	0.000225

Mini-Exercise – Operations Practice

Using the existing DataFrame (Day, Return, Cumulative Return):

1. Add a column "Absolute Return" → absolute value of daily return.
2. Add a column "Rolling volatility" → 3-day rolling standard deviation.
3. Add a column "Return Cubed" → cube of daily returns using `apply()`.
4. Compute mean, sum, min, max, std of "Return".

💡 Hint:

- "Absolute values: `abs()`"
- "Rolling std: `.rolling(window=3).std()`"
- "Cube: `apply(lambda x: x**3)`"

```
import pandas as pd

# Assume df is the existing DataFrame with Day, Return, Cumulative Return

# 1. Add Absolute Return column
df["Absolute Return"] = df["Return"].abs()

# 2. Add Rolling Volatility (3-day window)
df["Rolling Volatility"] = df["Return"].rolling(window=3).std()

# 3. Add Return Cubed column using apply()
df["Return Cubed"] = df["Return"].apply(lambda x: x**3)

# 4. Compute basic statistics on Return
mean_return = df["Return"].mean()
sum_return = df["Return"].sum()
min_return = df["Return"].min()
max_return = df["Return"].max()
std_return = df["Return"].std()

# Display DataFrame
print("Updated DataFrame:")
print(df)

# Display statistics
print("\nStatistics on Return:")
print(f"Mean Return: {mean_return}")
print(f"Sum of Returns: {sum_return}")
print(f"Min Return: {min_return}")
print(f"Max Return: {max_return}")
print(f"Std of Return: {std_return}")
```

Output

Updated DataFrame:

pgsql



	Day	Return	Cumulative Return	Absolute Return	Rolling Volatility	Return Cubed
0	1	0.010	1.010000	0.010	NaN	0.000001
1	2	-0.005	1.004950	0.005	NaN	-0.000000125
2	3	0.020	1.025049	0.020	0.012583	0.000008
3	4	-0.010	1.014798	0.010	0.015275	-0.000001
4	5	0.015	1.030020	0.015	0.015275	0.000003375

Statistics on Return:

mathematica



Mean Return: 0.006

Sum of Returns: 0.03

Min Return: -0.01

Max Return: 0.02

Std of Return: 0.01224744871391589

Topic 5: Handling Missing Data & Cleaning (Using Existing DataFrame)

1. Display Existing DataFrame (Context Reminder)

python

 Copier le code

```
import pandas as pd
import numpy as np


# Assume 'df' is the existing DataFrame with all columns
print("Existing DataFrame:")
print(df)
```

Example Output:

pgsql

 Copier le code

	Day	Return	Cumulative Return	Absolute Return	Rolling Volatility	Squared Return	Return Cub
0	1	0.010	1.010000	0.010	NaN	0.000100	0.000001
1	2	-0.005	1.004950	0.005	NaN	0.000025	-0.000000
2	3	0.020	1.025049	0.020	0.012583	0.000400	0.000008
3	4	-0.010	1.014798	0.010	0.015275	0.000100	-0.000001
4	5	0.015	1.030020	0.015	0.015275	0.000225	0.000003

 Note: The `Rolling Volatility` column contains `NaN` in the first rows due to the 3-day rolling window.

2. Detect Missing Values

python

 Copier le code

```
# Detect NaN
print("\nDetect missing values (bool):")
print(df.isna())

# Count missing values per column
print("\nCount of missing values per column:")
print(df.isna().sum())
```

Output:

mathematica

 Copier le code

```
Detect missing values (bool):
   Day  Return  Cumulative Return  Absolute Return  Rolling Volatility  Squared Return  Return C
0  False  False                False              False              True              False      Fal
1  False  False                False              False              True              False      Fal
2  False  False                False              False              False             False      Fal
3  False  False                False              False              False             False      Fal
4  False  False                False              False              False             False      Fal

Count of missing values per column:
Day                0
Return             0
Cumulative Return  0
Absolute Return    0
Rolling Volatility 2
Squared Return     0
Return Cubed       0
dtype: int64
```

3. Filling Missing Values

- Fill NaN in `Rolling Volatility` with `0` (or column mean if preferred):

python

 Copier le code

```
df["Rolling Volatility"] = df["Rolling Volatility"].fillna(0)
print("\nDataFrame after filling NaN in Rolling Volatility:")
print(df)
```

Output:

sql

 Copier le code

	Day	Return	Cumulative Return	Absolute Return	Rolling Volatility	Squared Return	Return Cut
0	1	0.010	1.010000	0.010	0.000	0.000100	0.000000
1	2	-0.005	1.004950	0.005	0.000	0.000025	-0.000000
2	3	0.020	1.025049	0.020	0.012583	0.000400	0.000000
3	4	-0.010	1.014798	0.010	0.015275	0.000100	-0.000001
4	5	0.015	1.030020	0.015	0.015275	0.000225	0.000003

4. Dropping Missing Values

- If there were NaN in critical columns (e.g., `Return` or `Cumulative Return`), drop them:

python

 Copier le code

```
df_clean = df.dropna(subset=["Return", "Cumulative Return"])
print("\nCleaned DataFrame (after dropping critical NaN):")
print(df_clean)
```

Ici, aucune ligne n'est supprimée car `Return` et `Cumulative Return` n'ont pas de NaN.

Mini-Exercise – Cleaning & Log Return Analysis

Context:

We have the existing DataFrame `df` with:

`Day`, `Return`, `Cumulative Return`, `Absolute Return`, `Rolling Volatility`, `Squared Return`, `Return Cubed`.

- Some values in `Rolling Volatility` are `NaN`.

Tasks:

1. Detect missing values per column.
2. Fill missing values in `Rolling Volatility` with the mean of the column.
3. Create a new column `Log Return` = `np.log(1 + Return)` using NumPy.
4. Identify the day with the maximum Log Return.
5. Compute mean and standard deviation of Log Return.
6. Identify the day with maximum Absolute Return.

💡 Hint:

- `~.isna()`, `.sum()` → detect missing values"
- `~.fillna(value=df["Rolling Volatility"].mean())` → fill NaN"
- `~np.log(1 + df["Return"])` → log return"
- `~.idxmax()` → find row of max value"
- `~.mean()`, `.std()` → statistics"

```
import numpy as np
import pandas as pd

# Assume 'df' is the existing DataFrame
print("Original DataFrame:")
print(df)

# 1. Detect missing values per column
print("\nMissing values per column:")
print(df.isna().sum())

# 2. Fill missing Rolling Volatility with column mean
rolling_mean = df["Rolling Volatility"].mean()
df["Rolling Volatility"] = df["Rolling Volatility"].fillna(rolling_mean)

# 3. Create Log Return column
df["Log Return"] = np.log(1 + df["Return"])

# 4. Identify the day with maximum Log Return
max_log_return_day = df["Log Return"].idxmax()

# 5. Compute mean and standard deviation of Log Return
mean_log_return = df["Log Return"].mean()
std_log_return = df["Log Return"].std()

# 6. Identify the day with maximum Absolute Return
max_abs_return_day = df["Absolute Return"].idxmax()

# 7. Display final DataFrame
print("\nCleaned DataFrame with Log Return:")
print(df)

# 8. Display results
print(f"\nDay with maximum Log Return: {df.loc[max_log_return_day, 'Day']}")
print(f"Mean of Log Return: {mean_log_return}")
print(f"Std of Log Return: {std_log_return}")
print(f"Day with maximum Absolute Return: {df.loc[max_abs_return_day, 'Day']}")
```

Example Output

Original DataFrame (with NaN in Rolling Volatility):

pgsql

 Copier le code

	Day	Return	Cumulative Return	Absolute Return	Rolling Volatility	Squared Return	Return Cubed
0	1	0.010	1.010000	0.010	NaN	0.000100	0.000001
1	2	-0.005	1.004950	0.005	NaN	0.000025	-0.000000
2	3	0.020	1.025049	0.020	0.012583	0.000400	0.000008
3	4	-0.010	1.014798	0.010	0.015275	0.000100	-0.000001
4	5	0.015	1.030020	0.015	0.015275	0.000225	0.000003

Missing values per column:

sql

 Copier le code

Day	0
Return	0
Cumulative Return	0
Absolute Return	0
Rolling Volatility	2
Squared Return	0
Return Cubed	0
dtype:	int64

Cleaned DataFrame with Log Return:

mathematica

Copier le code

	Day	Return	Cumulative Return	Absolute Return	Rolling Volatility	Squared Return	Return Cub
0	1	0.010	1.010000	0.010	0.014378	0.000100	0.000001
1	2	-0.005	1.004950	0.005	0.014378	0.000025	-0.000000
2	3	0.020	1.025049	0.020	0.012583	0.000400	0.000008
3	4	-0.010	1.014798	0.010	0.015275	0.000100	-0.000001
4	5	0.015	1.030020	0.015	0.015275	0.000225	0.000003

Results:

mathematica

Copier le code

Day with maximum Log Return: 3
Mean of Log Return: 0.009918
Std of Log Return: 0.011231
Day with maximum Absolute Return: 3

Topic 6: Grouping & Merge in Pandas

1. Grouping Data

- `groupby()` allows you to aggregate data by one or more columns.
- Common operations: `mean()`, `sum()`, `count()`, `std()`.
- Financial examples:
 - Average return per stock
 - Total volume traded per sector

```
import pandas as pd

# Sample DataFrame
data = {
    "Stock": ["A", "A", "A", "B", "B", "B"],
    "Day": [1, 2, 3, 1, 2, 3],
    "Return": [0.01, -0.005, 0.02, 0.015, -0.01, 0.005]
}
df = pd.DataFrame(data)

# Group by Stock and compute mean return
grouped = df.groupby("Stock")["Return"].mean()
print("Mean Return by Stock:")
print(grouped)
```

Output:

```

CSS

Stock
A    0.008333
B    0.003333
Name: Return, dtype: float64
```

Explanation:

- `df.groupby("Stock")` → splits the data into groups by Stock.
- `["Return"].mean()` → computes the mean of `Return` for each group.

Merge in Pandas – Detailed Explanation

1. What is Merge?

- `merge()` allows you to combine two DataFrames based on one or more columns (keys).
- Similar to SQL joins: inner, left, right, outer.
- Key idea: you define which column(s) to match and what kind of join you want.

2. Example 1: Inner Join

- Inner join keeps only rows that exist in both DataFrames.

python

```
import pandas as pd

# Prices DataFrame
df_prices = pd.DataFrame({
    "Stock": ["A", "B", "C"],
    "Price": [100, 200, 300]
})

# Sectors DataFrame
df_sector = pd.DataFrame({
    "Stock": ["A", "B", "D"],
    "Sector": ["Tech", "Finance", "Healthcare"]
})

# Inner merge
df_inner = pd.merge(df_prices, df_sector, on="Stock", how="inner")
print(df_inner)
```

Output:

css

	Stock	Price	Sector
0	A	100	Tech
1	B	200	Finance

Explanation:

- Only **A** and **B** exist in both DataFrames, so only these rows appear.
- **C** is in `df_prices` only → ignored.
- **D** is in `df_sector` only → ignored.

3. Example 2: Left Join

- Left join keeps all rows from the left DataFrame, missing matches from the right → NaN.

python



```
df_left = pd.merge(df_prices, df_sector, on="Stock", how="left")  
print(df_left)
```

Output:

css



	Stock	Price	Sector
0	A	100	Tech
1	B	200	Finance
2	C	300	NaN

Explanation:

- All rows from df_prices are kept.
- Stock C has no corresponding sector → Sector = NaN.

4. Example 3: Right Join

- Right join keeps all rows from the right DataFrame, missing matches from left → NaN.

python

```
df_right = pd.merge(df_prices, df_sector, on="Stock", how="right")
print(df_right)
```

Output:

css

	Stock	Price	Sector
0	A	100.0	Tech
1	B	200.0	Finance
2	D	NaN	Healthcare

Explanation:

- All rows from `df_sector` are kept.
- Stock `D` is not in `df_prices` → Price = NaN.

5. Example 4: Outer Join

- Outer join keeps all rows from both DataFrames, missing matches → NaN.

python

```
df_outer = pd.merge(df_prices, df_sector, on="Stock", how="outer")  
print(df_outer)
```

Output:

css

	Stock	Price	Sector
0	A	100.0	Tech
1	B	200.0	Finance
2	C	300.0	NaN
3	D	NaN	Healthcare

Explanation:

- Includes all stocks from both DataFrames.
- Missing values are filled with NaN.

7. Summary Table

Join Type	Keeps Rows From	Missing Matches
inner	Both DataFrames	Dropped
left	Left DataFrame	NaN
right	Right DataFrame	NaN
outer	Both DataFrames	NaN

Opening Files in Pandas – CSV & Excel

1. Reading a CSV File

- CSV = Comma-Separated Values, very common for datasets.
- Use `pd.read_csv()` to load into a DataFrame.

Example:

```
python

import pandas as pd

# Read CSV file
df_csv = pd.read_csv("data.csv")

# Show first 5 rows
print(df_csv.head())
```

Explanation:

- `"data.csv"` → path to your CSV file.
- `.head()` → displays the first 5 rows.
- Optional parameters:
 - `sep=";"` → if values are separated by `;` instead of `,`
 - `index_col=0` → set first column as index
 - `parse_dates=["Date"]` → automatically convert Date column to datetime

2. Reading an Excel File

- Excel files can have multiple sheets.
- Use `pd.read_excel()` to load.

Example:

```
python
```

```
# Read Excel file, sheet named "Prices"  
df_excel = pd.read_excel("data.xlsx", sheet_name="Prices")  
  
# Show first 5 rows  
print(df_excel.head())
```

Explanation:

- `"data.xlsx"` → path to your Excel file.
- `sheet_name="Prices"` → choose which sheet to read.
- Optional parameters:
 - `index_col=0` → set first column as index
 - `parse_dates=["Date"]` → convert Date column to datetime

Function	Description	Example Code	Example Output
<code>pd.read_csv()</code>	Read a CSV file into a DataFrame	<code>df = pd.read_csv("data.csv")</code>	<pre>Date Stock Return 0 2025-01-01 A 0.01 1 2025-01-02 A 0.02</pre>
<code>pd.read_excel()</code>	Read an Excel file	<code>df = pd.read_excel("data.xlsx", sheet_name="Prices")</code>	<pre>Date Stock Price 0 2025-01-01 A 100 1 2025-01-02 A 101</pre>
<code>df.head(n)</code>	Display first n rows	<code>df.head(3)</code>	<pre>Date Stock Return 0 2025-01-01 A 0.01 1 2025-01-02 A 0.02 2 2025-01-03 B -0.01</pre>
<code>df.tail(n)</code>	Display last n rows	<code>df.tail(2)</code>	<pre>Date Stock Return 3 2025-01-04 B 0.015 4 2025-01-05 A -0.005</pre>
<code>df.info()</code>	Summary of DataFrame	<code>df.info()</code>	<pre><class 'pandas.DataFrame'>... RangeIndex: 5 entries, 0 to 4</pre>
<code>df.describe()</code>	Summary statistics of numeric columns	<code>df.describe()</code>	<pre>Return count 5 mean 0.005 std 0.012 min -0.01 max 0.02</pre>
<code>df.isna()</code> / <code>df.isna().sum()</code>	Detect missing values	<code>df.isna().sum()</code>	<pre>Date 0 Stock 0 Return 1</pre>
<code>df.fillna(value)</code>	Fill missing values	<code>df["Return"].fillna(0, inplace=True)</code>	<pre>Return 0.01 0.02 0.00 0.015 -0.005</pre>

<code>df.groupby()</code>	Group data by column(s)	<code>df.groupby("Stock")</code> <code>["Return"].mean()</code>	Stock A 0.008333 B 0.003333
<code>df.merge()</code>	Merge two DataFrames	<code>pd.merge(df1, df2, on="Stock", how="left")</code>	Stock Price Sector A 100 Tech B 200 Finance C 300 NaN
<code>df.sort_values()</code>	Sort DataFrame by column(s)	<code>df.sort_values("Return", ascending=False)</code>	Stock Return 2 A 0.02 0 A 0.01
<code>df.set_index()</code>	Set a column as index	<code>df.set_index("Date", inplace=True)</code>	Index = Date column
<code>df.reset_index()</code>	Reset index to default	<code>df.reset_index(inplace=True)</code>	Index = 0,1,2,...
<code>df.apply()</code>	Apply a function to column/row	<code>df["Return"].apply(lambda x: x**2)</code>	Return 0 0.0001 1 0.0004 2 0.000225
<code>df.loc[]</code> / <code>df.iloc[]</code>	Indexing by label / position	<code>df.loc[0, "Return"]</code> → 0.01	0.01
<code>df.cumsum()</code>	Cumulative sum	<code>df["Return"].cumsum()</code>	0 0.01 1 0.03 2 0.02
<code>df.shift()</code>	Shift data for lag/lead	<code>df["Return"].shift(1)</code>	0 NaN 1 0.01 2 0.02
<code>df.resample()</code>	Resample time series	<code>df.resample("W").mean()</code>	Weekly average return

Exercise – Portfolio Statistical Analysis with Equal Weights

Dataset:

Date , CAC 40 , DAX , FTSE , S&P , IBOVESPA , IGPA , KOSPI , US T-Bill , SSE

Tasks

Step 1: Load & Prepare Data

1. Load the CSV/Excel file into a Pandas DataFrame.
2. Convert the Date column to datetime format.
3. Set Date as the index of the DataFrame.
4. Ensure all numeric columns are floats (replace commas , with dots . if needed).
5. Fill any missing values with the mean of the corresponding column.

Step 2: Compute Log Returns

1. Calculate daily log returns for each index:

$$r_t = \ln \left(\frac{P_t}{P_{t-1}} \right)$$

2. Calculate cumulative log returns for each index:

$$R_{cum} = \sum r_t$$

💡 Hint: Use np.log() and .shift(1)

Step 3: Asset Statistics

1. For each index, compute:
 - Mean daily log return
 - Standard deviation (volatility)
 - Minimum daily return
 - Maximum daily return
 - Sharpe ratio (assume risk-free rate = 0):

$$\text{Sharpe} = \frac{\text{Mean}}{\text{Std}}$$

2. Identify the most attractive asset in terms of risk-adjusted return (highest Sharpe ratio).
3. Identify the most volatile asset (highest standard deviation).

Step 4: Portfolio with Equal Weights

1. Assume a portfolio including all indices with equal weights.
2. Compute the portfolio daily log return:

$$R_{p,t} = \sum_i w_i r_{i,t}$$

3. Compute the cumulative portfolio log return.
 4. Compute portfolio statistics: mean, standard deviation, min, max, Sharpe ratio.
-

Step 5: Insights

- Compare individual assets' Sharpe ratios with the portfolio Sharpe ratio.
- Discuss the diversification benefits of combining multiple assets.
- Comment on which indices contribute most to portfolio volatility.

💡 Hint: Use `.mean()`, `.std()`, `.min()`, `.max()`, `.cumsum()`, and `np.dot(weights, log_returns.T)`

Solution – Portfolio Statistical Analysis with Equal Weights

python

 Copier le code

```
import pandas as pd
import numpy as np

# --- Step 1: Load & Prepare Data ---
df = pd.read_csv("indices.csv", parse_dates=["Date"], dayfirst=True)
df.set_index("Date", inplace=True)

# Convert numeric columns to float
df = df.apply(lambda x: x.str.replace(',', '.').astype(float) if x.dtype=='object' else x)

# Fill missing values with column mean
df.fillna(df.mean(), inplace=True)

print("DataFrame head:\n", df.head())
```

Output (simulated):

yaml

 Copier le code

	CAC 40	DAX	FTSE	S&P	IBOVESPA	IGPA	KOSPI	US T-Bill	SSE
Date									
2016-10-28	4548.58	10526.16	6845.4	2088.66	64158	20792.95	1974.40	216.40	3104.27
2016-10-31	4509.26	10370.93	6790.5	2085.18	63258	20587.41	1967.53	216.40	3100.49
2016-11-01	4470.28	10325.88	6693.3	2131.52	61201	20801.85	1979.65	216.39	3122.44
2016-11-02	4414.67	10259.13	6806.9	2139.56	59184	20915.76	1980.55	216.39	3102.73
2016-11-03	4411.68	10456.95	6843.1	2163.26	59657	20939.45	1974.58	216.39	3128.94

Step 2: Daily Log Returns

python

 Copier le code

```
log_returns = np.log(df / df.shift(1))
log_returns = log_returns.dropna()

cum_log_returns = log_returns.cumsum()

print("Daily Log Returns:\n", log_returns.head())
print("\nCummulative Log Returns:\n", cum_log_returns.head())
```

Output (simulated):

yaml

 Copier le code

Daily Log Returns:									
	CAC 40	DAX	FTSE	S&P	IBOVESPA	IGPA	KOSPI	US T-Bill	
Date									
2016-10-31	-0.00864	-0.01468	-0.00806	-0.00172	-0.01405	-0.00992	-0.00348	0.00000	-0.00123
2016-11-01	-0.00865	-0.00432	-0.01478	0.02203	-0.03333	0.00687	0.00613	-0.00005	0.00715
2016-11-02	-0.01295	-0.00648	0.01691	0.00377	-0.03454	0.00545	0.00045	-0.00005	-0.00638
2016-11-03	-0.00068	0.01915	0.00543	0.00355	0.00804	0.00114	-0.00298	0.00000	0.00850
2016-11-04	-0.00769	0.00245	0.01001	0.00191	0.01809	-0.00018	-0.00440	0.00005	-0.00119
Cumulative Log Returns:									
	CAC 40	DAX	FTSE	S&P	IBOVESPA	IGPA	KOSPI	US T-Bill	
Date									
2016-10-31	-0.00864	-0.01468	-0.00806	-0.00172	-0.01405	-0.00992	-0.00348	0.00000	-0.00123
2016-11-01	-0.01729	-0.01899	-0.02284	0.02031	-0.04738	-0.00305	0.00265	-0.00005	0.00592
2016-11-02	-0.03024	-0.02547	-0.00593	0.02408	-0.08191	0.00240	0.00310	-0.00010	-0.00046
2016-11-03	-0.03092	-0.00632	0.00050	0.02763	-0.07387	0.00354	0.00012	0.00010	0.00804
2016-11-04	-0.03861	-0.00387	0.01051	0.02954	-0.05578	0.00336	-0.00428	0.00015	0.00685

Step 3: Asset Statistics

python

```
stats_assets = pd.DataFrame(index=df.columns)
stats_assets["Mean"] = log_returns.mean()
stats_assets["Std"] = log_returns.std()
stats_assets["Min"] = log_returns.min()
stats_assets["Max"] = log_returns.max()
stats_assets["Sharpe"] = stats_assets["Mean"] / stats_assets["Std"]

print("\nStatistics for Each Asset:\n", stats_assets)
```

Output (simulated):

mathematica

	Mean	Std	Min	Max	Sharpe
CAC 40	-0.00345	0.0123	-0.0256	0.0187	-0.280
DAX	-0.00175	0.0118	-0.0221	0.0192	-0.148
FTSE	-0.00095	0.0105	-0.0182	0.0168	-0.090
S&P	0.00214	0.0098	-0.0153	0.0203	0.218
IBOVESPA	-0.00223	0.0182	-0.0353	0.0301	-0.123
IGPA	0.00105	0.0109	-0.0186	0.0200	0.096
KOSPI	0.00042	0.0112	-0.0195	0.0184	0.037
US T-Bill	0.00000	0.0000	0.0000	0.0000	NaN
SSE	0.00145	0.0120	-0.0205	0.0221	0.121

Step 4: Equal-Weight Portfolio

python

```
n_assets = len(log_returns.columns)
weights = np.array([1/n_assets]*n_assets)

portfolio_return = log_returns.dot(weights)

portfolio_stats = {
    "Mean": portfolio_return.mean(),
    "Std": portfolio_return.std(),
    "Min": portfolio_return.min(),
    "Max": portfolio_return.max(),
    "Sharpe": portfolio_return.mean() / portfolio_return.std()
}

print("\nPortfolio Statistics (Equal Weights):\n", portfolio_stats)
```

Output (simulated):

csharp

```
Portfolio Statistics (Equal Weights):
{'Mean': 0.00054, 'Std': 0.0085, 'Min': -0.0182, 'Max': 0.0167, 'Sharpe': 0.064}
```